

NodeJs 应用 性能分析优化 & 分布式设计

提纲

- 从实例开始
- 性能分析&优化方法
 - 资源占用分析
 - CPU、内存
 - 文件IO、网络IO
 - 慢代码分析
 - V8 掠影
 - 内存、堆栈、GC、预编译
- 分布式设计
 - 单机
 - 集群

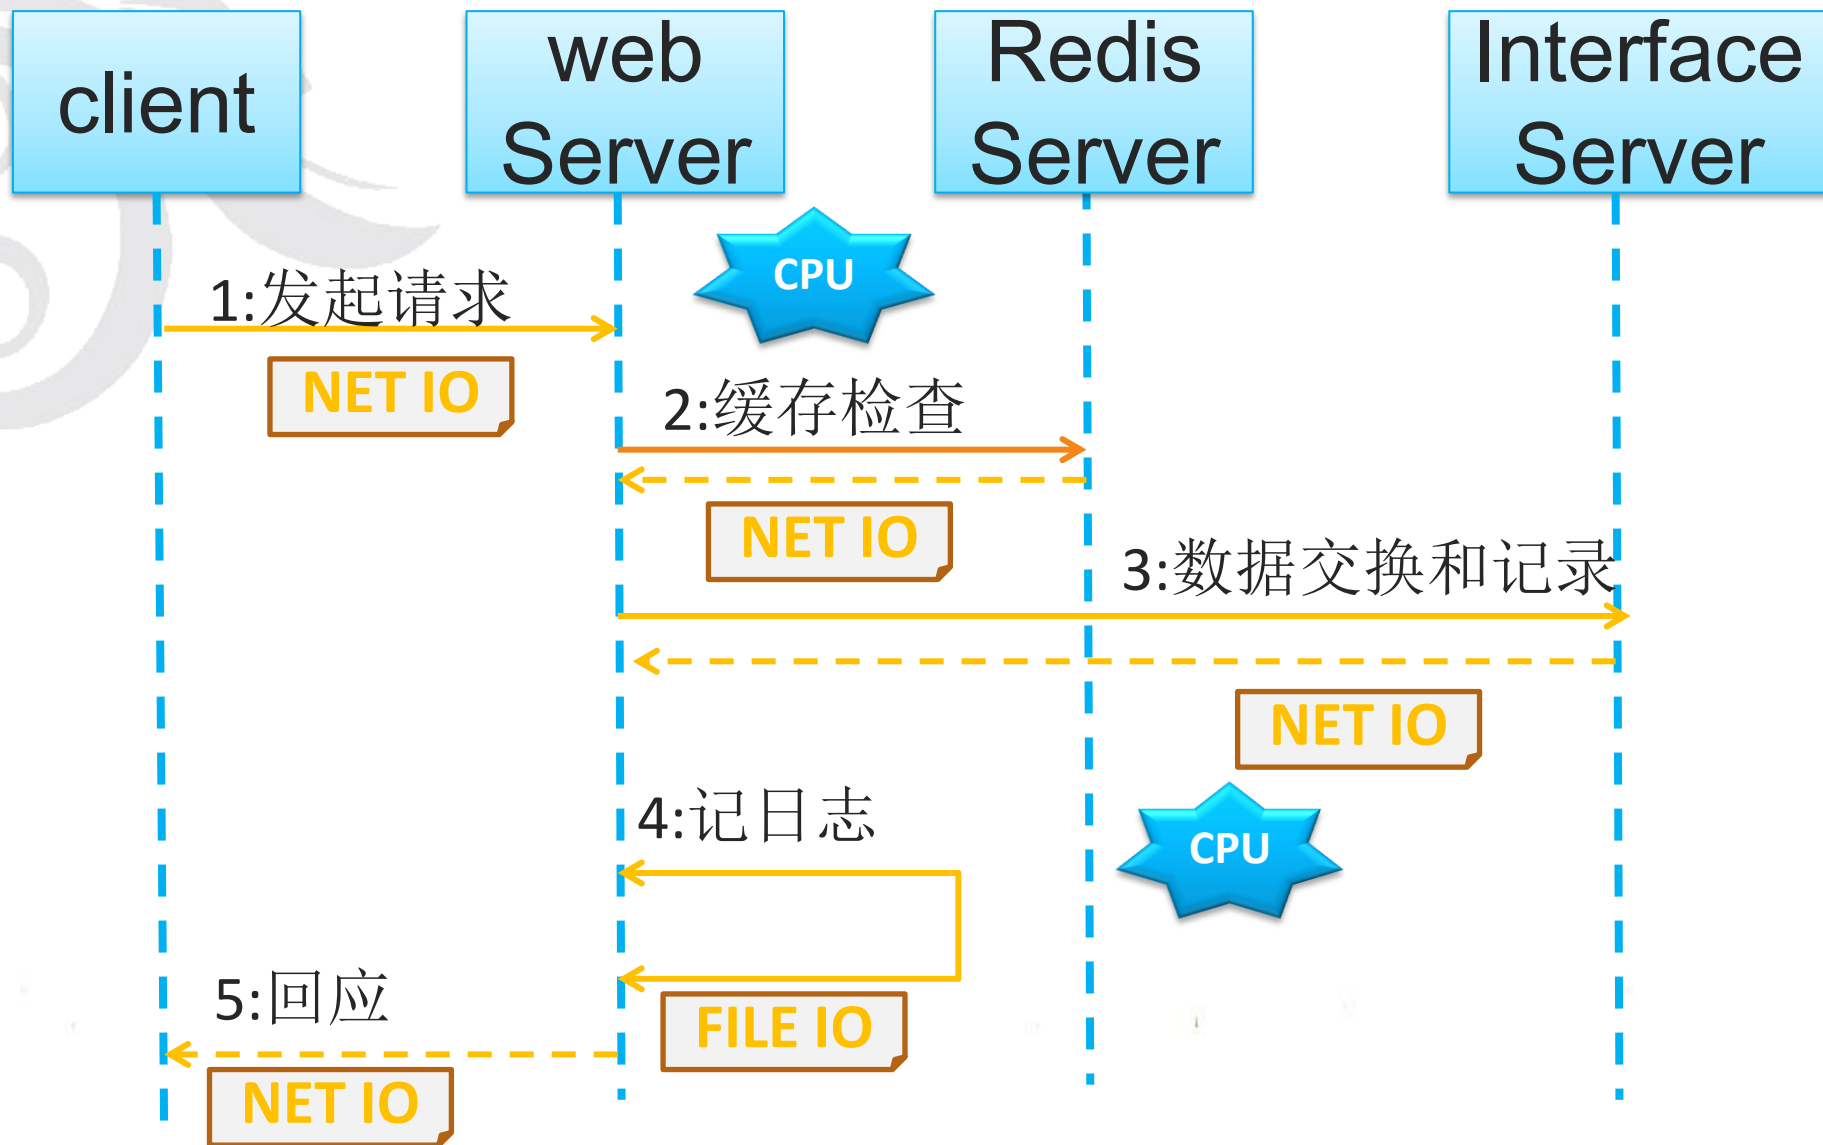




性能分析和优化

性能分析流程、工具、方法、优化要点

从实例开始



工具

- Linux tools
 - pidstat/iostat/vmstat
 - sar/top/lsof
- node lib
 - [v8-profiler](#)
 - [Benchmark.js](#)
- V8 tools
 - node-v0.6.2/deps/v8/tools
 - linux-tick-processor
 - ll_prof.py
 - run-valgrind.py

CPU 占用资源分析

利用率：

用户进程/内核/中断/IO等待/空闲

us / sy/(hi/si)/wa /id →top

建议值：

usr/sys:65%-75% / 30%-35%

分析:

top (1->shift+h)

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1651	lizhe.zc	20	0	647m	47m	5272	S	24.9	2.4	0:44.63	node
1652	lizhe.zc	20	0	647m	47m	5272	S	3.3	2.4	0:05.73	SignalSender
1653	lizhe.zc	20	0	647m	47m	5272	S	0.3	2.4	0:00.93	node
1654	lizhe.zc	20	0	647m	47m	5272	S	0.7	2.4	0:01.42	node
1655	lizhe.zc	20	0	647m	47m	5272	S	0.0	2.4	0:01.16	node
1656	lizhe.zc	20	0	647m	47m	5272	S	1.3	2.4	0:00.84	node

CPU 占用资源分析

- `pidstat -p 1651 -t 1 100`

node 应用

| -node 主线程

| -SignalSender线程 profile sampling

| -4个libuv线程 iowatcher etc

TGID	TID	%usr	%system	%guest	%CPU	CPU	Command
1651	-	3.92	4.90	0.00	8.82	0	node
-	1651	3.92	4.90	0.00	8.82	0	__node
-	1652	0.98	0.00	0.00	0.98	1	__SignalSender
-	1653	0.00	0.00	0.00	0.00	0	__node
-	1654	0.00	0.00	0.00	0.00	1	__node
-	1655	0.00	0.00	0.00	0.00	1	__node
-	1656	0.00	0.00	0.00	0.00	0	__node

CPU 优化/利用点

- 代码适应 V8
- 减少GC
- 多进程
- 原生代码
- 模板选型
- 复杂计算业务逻辑转移
 - Java/C app
 - Gearman类 任务分发 (异步化)
 - MQ
- 语言层面
 - eval
 - setInterval/setTimeout (arg)
 - Primitive operations
 - Regexp
 - async+parallel
 - Object constructor
 - Array Pre-allocate
 - 精简解析 (httpparser)

文件IO 占用资源分析

- `pidstat -d -p 1651 -t 1 100`
 - `kB_rd/s` `kB_wr/s` `kB_ccwr/s`
- `iostat -x vda2 3 5`
 - `%util`
 - `await`
- `sar -b`
 - `rtps / wtps / (bread/s) / (bwrtn/s)`
- 优化点
 - IO分散
 - stream读取大文件
 - async
 - unwatcher

网络IO占用分析

- `sar -n DEV 1 10`
 - `IFACE rxpck/s txpck/s rxkB/s`
- `sar -n SOCK 1 10`
 - `tcpsck udpsck`
- `sar -n TCP 1 10`
 - `iseg/s oseg/s`
- 优化点：
 - `maxsockets`
 - `timeout`
 - `response header (expire ..)`
 - `request (no cookies..)`
 - `pool`
 - `sync/async (getaddrinfo/ gethostbyname / ares_gethostbyname)`
 - 分段读取
 - 压缩传输 (`msgpack/bin/gzip`)

内存占用

- free/vmstat
 - Cached/buffered/swpd
- sar -B 1 5
 - (pgpgin/s) / (pgpgout/s) / (pgscank/s)
 - (pgscand/s) / (pgsteal/s) / (%vmeff)
- sar -r 1 5
 - Kbmemfree+kbbuffers+kbcached
- pidstat -r -p 1813 1 10
 - minflt/s majflt/s VSZ RSS
- pidstat -s -p 1813 1 10
 - minflt/s majflt/s VSZ RSS

内存占用

- 优化点：
 - 整体
 - 加入 Buffer(堆外内存)
 - 加大最大内存设置
 - `--max_old_space_size = 1900` (64bit os)
 - `--stack_size=4096`
 - `--max_new_space_size=10000`
 - `--use_big_map_space` (慎用)
 - 语言层面
 - 局部变量
 - `Try {bigset}catch ()`
 - `> try {fn}`
 - <http://jsperf.com/try-catch-performance-overhead>
 - `TypedArray`
 - `Cache`
 - `With`
 - 对象转换、`copy`
 - `String concat ...`

开始说说 代码性能+V8

- Benchmark 测试
 - 单元测试不仅仅只验证**正确性**

```
var suite = new Benchmark.Suite;  
// add tests  
suite.add( 'RegExp#test' , function() {  
    /o/.test( 'Hello World!' );  
})  
.add( 'String#match' , function() {  
    !! 'Hello World!' .match(/o/);  
})  
// add listeners  
.on( 'cycle' , function(event, bench) {  
    ...  
})  
.on( 'complete' , function() {  
    ...  
})  
.run({ 'async' : true });
```

Sample

data length

\u0000

data bytes

6

\u0000

T

A

O

b

A

O

Step 1

```
Parser.prototype.parse1 = function (s) {  
  var l = '';  
  for (var i = 0; i < s.length; i++) {  
    if (s[i] == '\u0000') {  
      l = Number(l);  
      this.emit('data', s.substr(i + 1, l));  
      return this.parse1(s.substr(i + 1 + l));  
    } else {  
      l += s[i];  
    }  
  }  
  return s;  
};
```

Step 1-Stress

```
var p = new Parser();  
var NOF_RUNS = 1000;  
var start = Date.now();  
for (var j = 0; j < RUN_NUMBERS; j++) {  
    p.parse3(fakeInput);  
}  
var end = Date.now();  
  
var timeSpent = end - start;  
console.log(timeSpent + ' ms');
```

400 ms

Step 1—key profile

[JavaScript]:

ticks	total	nonlib	name
-------	-------	--------	------

38	15.8%	21.0%	Stub: SubStringStub
----	-------	-------	---------------------

2	0.8%	1.1%	Stub: StringAddStub
---	------	------	---------------------

2	0.8%	1.1%	LazyCompile: *Parser.parse1 /work/project/stress/src/BinFile.js:10
---	------	------	--------------------------------------------------------------------

1	0.4%	0.6%	Stub: StringAddStub {1}
---	------	------	-------------------------

1	0.4%	0.6%	LazyCompile: *substr native string.js:698
---	------	------	-------------------------------------------

[GC]:

ticks	total	nonlib	name
-------	-------	--------	------

151	62.9%		
-----	-------	--	--

GC成本随长时间存活
对象的个数线性上涨

pause=9 mutator=7 gc=s external=0 mark=0 sweep=0 sweepns=0 compact=0
total_size_before=22049776 total_size_after=22001000 holes_size_before=335256
holes_size_after=335256 allocated=16776520 promoted=7174952

Memory allocator, used: 84180992, available: 1450934272

New space, used: 9551576, available: 7225640

Old pointers, used: 605384, available: 1710936, waste: 160

Old data space, used: 203208, available: 300712, waste: 16

Code space, used: 361472, available: 126208, waste: 0

Map space, used: 39704, available: 207624, waste: 4640

Cell space, used: 8128, available: 251968, waste: 0

Large object space, used: 13025280, available: 1450926016

Step 2

```
Parser.prototype.parse1 = function (s) {  
  var l = '';  
  for (var i = 0; i < s.length; i++) {  
    if (s[i] == '\u0000') {  
      l = Number(l);  
      this.emit('data', s.substr(i + 1, l));  
      return this.parse1(s.substr(i + 1 + l));  
    } else {  
      l += s[i];  
    }  
  }  
  return s;  
};
```

Step 2

```
Parser.prototype.parse1 = function (s) {  
  var l = '';  
  for (var i = 0; i < s.length; i++) {  
    if (s[i] == '\u0000') {  
      l = Number(l);  
      this.emit('data', s.substr(i + 1, l));  
      s = s.substr(i + 1 + l);  
      i = 0;  
      l = '';  
    } else {  
      l += s[i];  
    }  
  }  
  return s;  
};
```

170 ms

Step2 -profile

[JavaScript]:

	ticks	total	nonlib	name
--	-------	-------	--------	------

42	18.8%	44.2%	Stub: SubStringStub
----	-------	-------	---------------------

3	1.3%	3.2%	Stub: StringAddStub
---	------	------	---------------------

2	0.9%	2.1%	LazyCompile: *Parser.parse2 /mnt/share/stress/src/BinFile.js:25
---	------	------	-----------------------------------------------------------------

1	0.4%	1.1%	LazyCompile: b native v8natives.js:1264
---	------	------	-----------------------------------------

[GC]:

	ticks	total	nonlib	name
--	-------	-------	--------	------

36	16.1%			
----	-------	--	--	--

pause=0 mutator=1 gc=s external=0 mark=0 sweep=0 sweepns=0 compact=0
total_size_before=7550080 total_size_after=3394272 holes_size_before=69824
holes_size_after=69824 allocated=4148080 promoted=0

Memory allocator, used: 71888896, available: 1463226368

New space, used: 22560, available: 4171744

Old pointers, used: 2060512, available: 245784, waste: 2056

Old data space, used: 252568, available: 259256, waste: 240

Code space, used: 415616, available: 88320, waste: 0

Map space, used: 39704, available: 215752, waste: 4640

Cell space, used: 8128, available: 251968, waste: 0

Large object space, used: 724992, available: 1463218112

Step 3

```
Parser.prototype.parse3 = function (s) {  
  var l = '';  
  //方法3  
  var j = 0;  
  
  for (var i = 0; i < s.length; i++) {  
    if (s[i] == '\u0000') {  
      l = Number(l);  
      this.emit('data', s.substr(i + 1, l));  
      i += l;  
      j = i + 1;  
    } else {  
      l += s[i];  
    }  
  }  
  return s.substr(j);  
};
```

11 ms

Step3-profile

[JavaScript]:

ticks	total	nonlib	name
-------	-------	--------	------

1	0.7%	3.4%	Stub: CallFunctionStub
---	------	------	------------------------

[GC]:

ticks	total	nonlib	name
-------	-------	--------	------

20	13.8%		
----	-------	--	--

pause=1 mutator=2 gc=s external=0 mark=0 sweep=0 sweepns=0
compact=0 total_size_before=2880944 total_size_after=2766424
holes_size_before=18528 holes_size_after=30208 allocated=790456
promoted=671920

Memory allocator, used: 70520832, available: 1464594432

New space, used: 262136, available: 786440

Old pointers, used: 1232664, available: 34592, waste: 712

Old data space, used: 231136, available: 28936, waste: 24

Code space, used: 434112, available: 53568, waste: 0

Map space, used: 54208, available: 193120, waste: 4640

Cell space, used: 8624, available: 243344, waste: 0

Large object space, used: 667648, available: 1464586176

Step4

```
Parser.prototype.parse4 = function (s) {
```

```
  var l = 0, i = 0;
```

```
  while (i < s.length) {
```

```
    var ch = s.charCodeAt(i);
```

```
    if (ch === 0) {
```

```
      this.emit('data', s.substr(i + 1, l));
```

```
      i += l + 1;
```

```
      l = 0;
```

```
    } else {
```

```
      l = l * 10 + ch;
```

```
      i ++;
```

```
    }
```

```
  }
```

```
};
```



50X

8 ms

Step4-profile

[JavaScript]:

ticks	total	nonlib	name
-------	-------	--------	------

[GC]:

ticks	total	nonlib	name
-------	-------	--------	------

17 12.4%

pause=3 mutator=1 gc=s external=0 mark=0 sweep=0 sweepns=0
compact=0 total_size_before=2889880 total_size_after=2769744
holes_size_before=20472 holes_size_after=25392 allocated=786184
promoted=666304

Memory allocator, used: 70520832, available: 1464594432

New space, used: 262136, available: 786440

Old pointers, used: 1231248, available: 36048, waste: 672

Old data space, used: 231080, available: 20880, waste: 8

Code space, used: 438976, available: 64960, waste: 0

Map space, used: 54152, available: 201304, waste: 4640

Cell space, used: 8608, available: 243360, waste: 0

Large object space, used: 667648, available: 1464586176

预编译和v8代码优化 日志

[**optimizing**: Queue.push / 25d70710ba79 - took 0.064 ms]

Compiled: 33 functions with 37333 byte source size in 31.198000ms.

[marking NonStringToString 0xc69df07d020 for recompilation]

Bailout in HGraphBuilder: @"NonStringToString": call to a JavaScript runtime function

[disabled optimization for: NonStringToString / c69df07d021]

[marking Buffer.write 0x143784371b80 for recompilation]

Bailout in HGraphBuilder: @"Buffer.write": SwitchStatement: non-literal switch label

Nodejs prof分析方法

1.Linux perf + node deep prof

```
perf record -R -e cycles -c 10000 -f node ../script.js --ll-prof  
ll_prof.py --disasm-top=10
```

2.Node parameter

Optimization:

- trace_opt (trace lazy optimization)
- trace_opt_stats (trace lazy optimization statistics)
- trace_deopt (trace deoptimization)
- trace_bailout (print reasons for falling back to using the classic V8 backend)

GC:

- trace_gc (print one trace line following each garbage collection)
- trace_gc_nvp (print one detailed trace line in name=value format after each garbage collection)
- print_cumulative_gc_stat (print cumulative GC statistics in name=value format on exit)
- trace_gc_verbose (print more details following each garbage collection)

3.Manual

--noprof-auto

profiler.startProfiling('startup'); - start/resume collection of data

profiler.stopProfiling - pause collection of data

v8-profiler

- ```
var profiler = require('v8-profiler');
profiler.startProfiling('startup');
slowStartupFoo();
profiler.stopProfiling('startup');
profiler.takeSnapshot('beforeLeak');
leakyFoo();
profiler.takeSnapshot('afterLeak');
```

# Node App 应用层面运维建议


- 定期收集运行信息（建议秒级别）
  - process.memoryUsage()
    - { rss, heapTotal, heapUsed }
  - process.cpuUsage()
    - eio\_init、 req\_init、 handle\_init,
    - stream\_init、 tcp\_init、 udp\_init,
    - prepare\_init、 check\_init
    - idle\_init、 async\_init
    - timer\_init:、 process\_init、 fs\_event\_init
- 定期开启profiler
  - 收集关键函数调用时间
  - 收集堆栈信息
- 其它IO收集
  - 请求数、响应时间
  - 内部系统交互响应时间等



## 分布式设计探讨

# 分布式设计（探讨）

- 单机：
  - 多进程（domain socket）
    - cluster
    - multi-node
- 集群
  - 节点无交互
    - Proxy（nginx proxy..）
    - LVS..
  - 节点有交互
    - RPC（缺点？）
      - thrift、rest、web services
- 高并发系统特性
  - 消息交互
  - 无状态
  - 异步？



# Nodejs集群（复杂计算逻辑 + 异构系统）

- ZEROMQ
  - 跨多种传输协议和方式
    - 进程内通讯
    - **IPC**
    - **TCP**
    - 广播
  - 多连接模型
    - REQ/REP
    - PUB/SUB
    - PUSH/PULL
  - 全局拓扑
    - 智能感知路由
  - 无锁
  - 异步消息交互
  - 低延迟高并发
  - 接口高度一致



# REQ/REP模型




Figure 39 – Basic request-reply



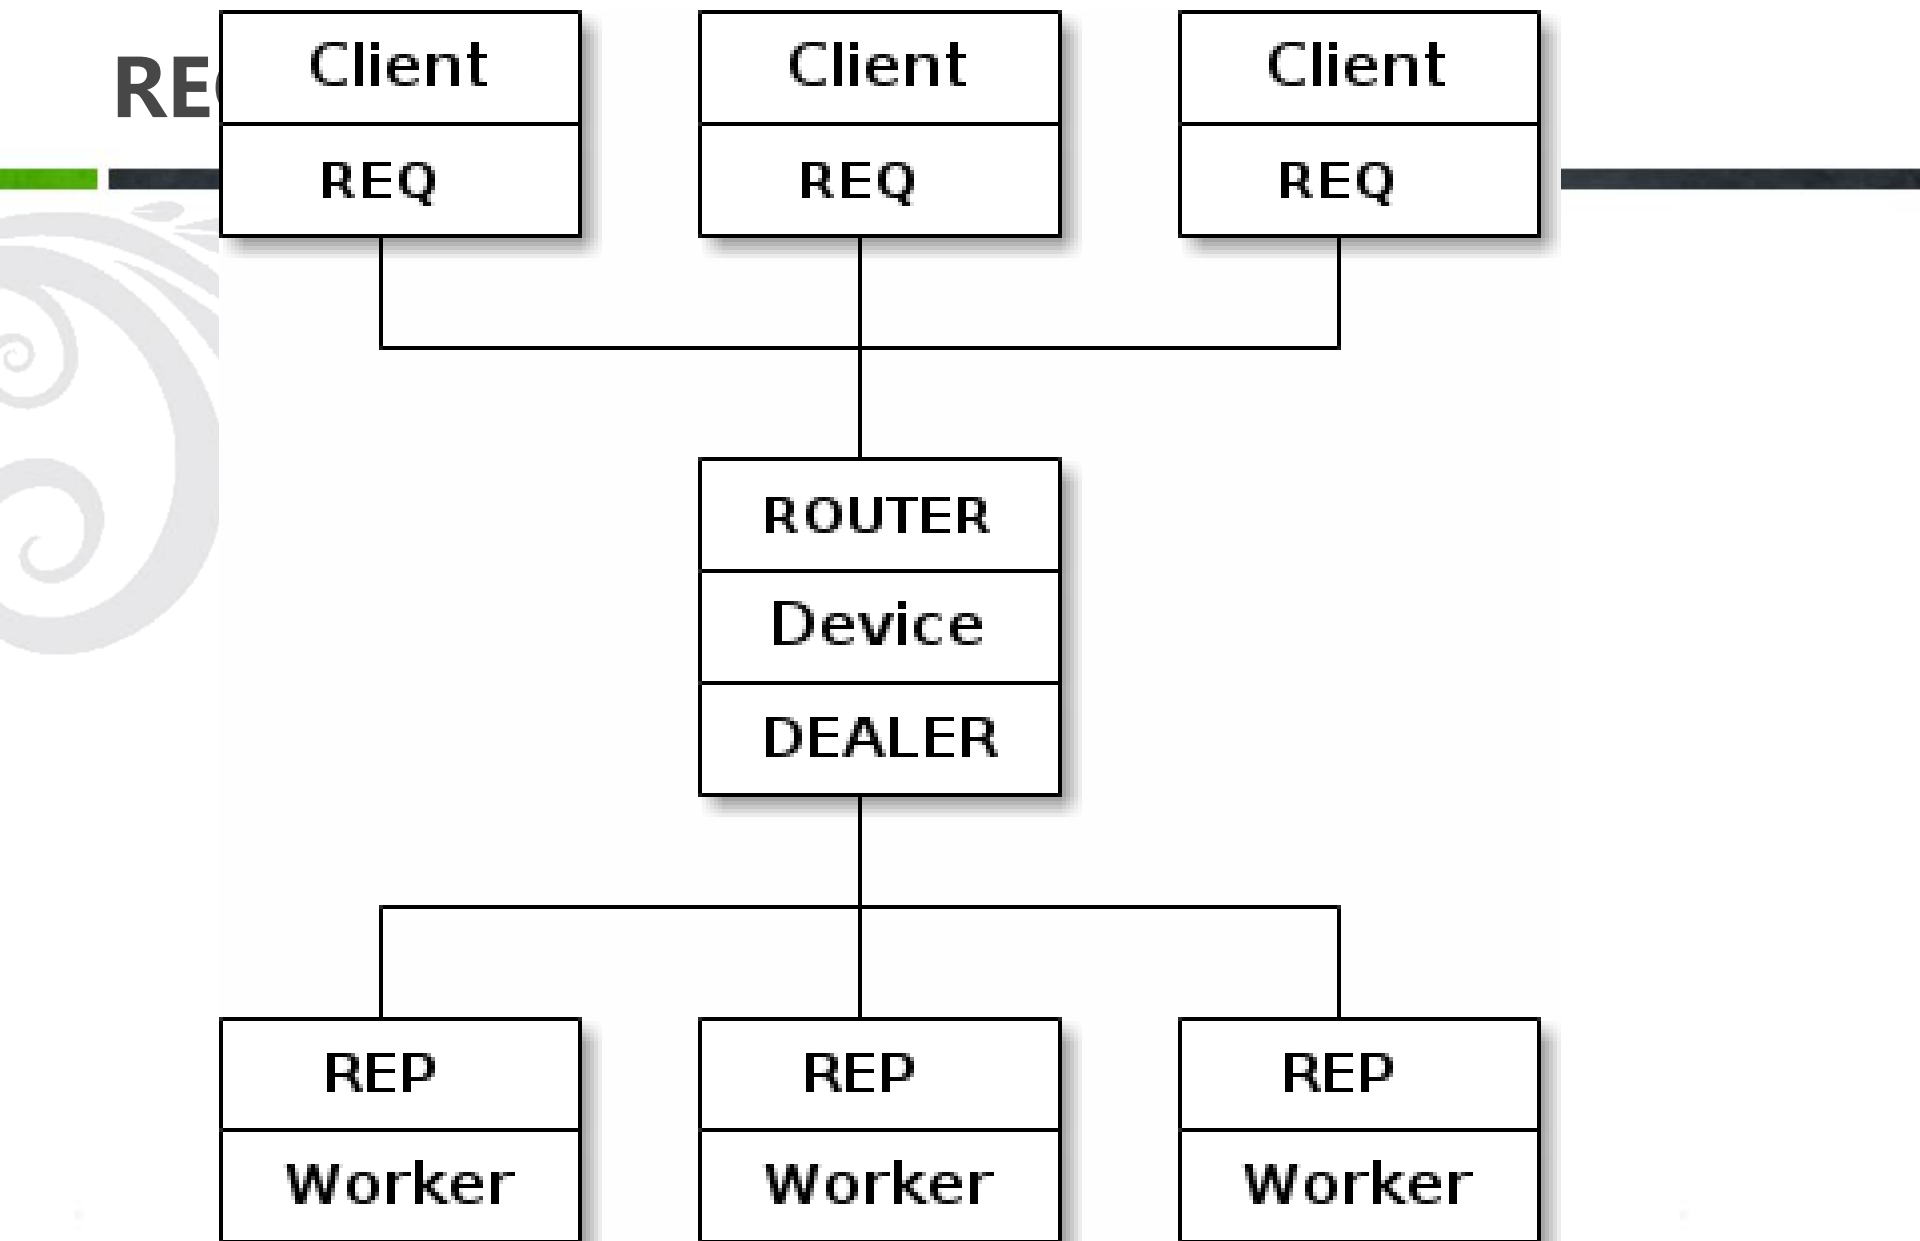


Figure 40 – Stretched request-reply

# REQ/REP模型




Figure 41 – Stretched request-reply with LRU

# REQ/REP模型

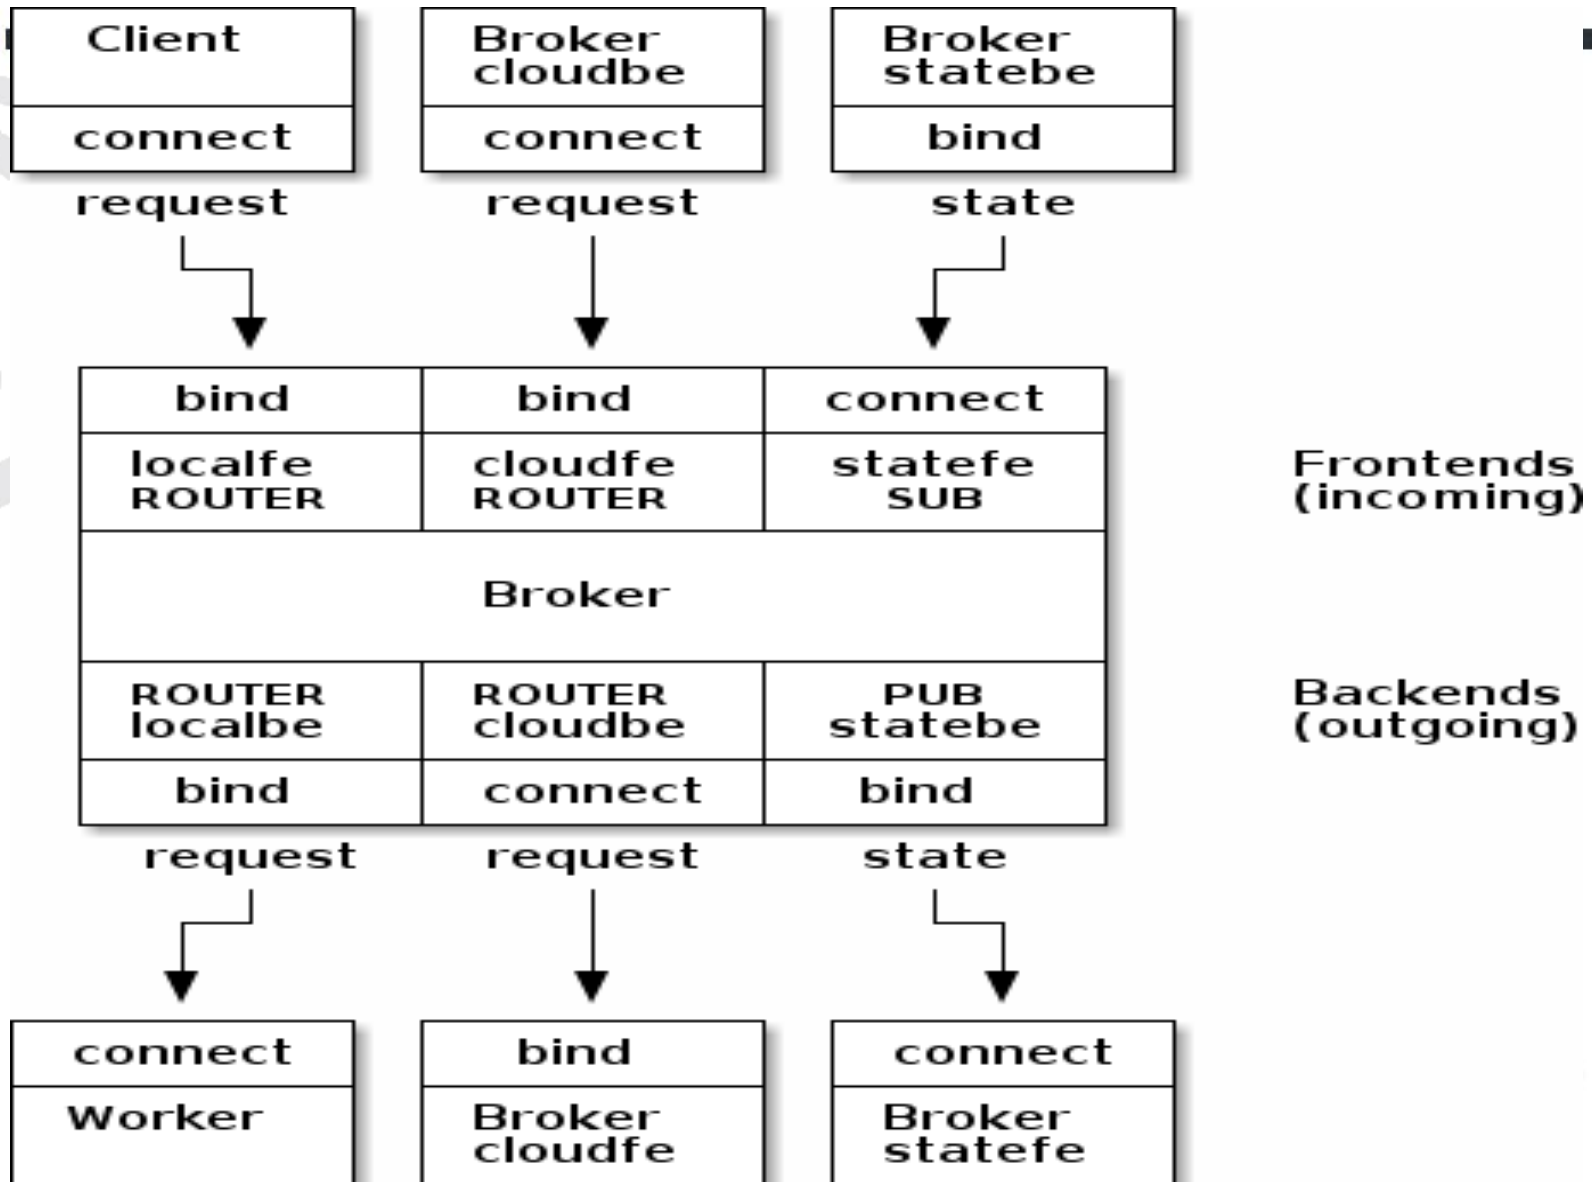


Figure 53 — Broker socket arrangement

# REQ/REP模型




Figure 52 – Cross connected brokers in federation model

# Sample

```
var zmq = require('zmq')
 , sock = zmq.socket('rep');
var i = 0;
sock.bindSync(url);
sock.on('message', function(msg){
});
```

代码不用变

url :

‘ipc:///tmp/zmq’ -- 进程间通讯  
‘tcp://\*:23456’ -- 网络



# 推荐：

- 编程规范：
  - <http://cnodejs.org/blog/?p=4739>
  - <https://github.com/windyrobin/iFrame/>
- [Blazing fast node.js: 10 performance tips from LinkedIn Mobile](#)
- [Efficient JavaScript](#)
- [JavaScript performance playground](#)



# Q & A

